

# Image Captioning using Attention Mechanism

## Week 2: Introduction to Convolutional Neural Networks (CNNs)



Welcome to the **second week** of our project course! This week, we will dive into Convolutional Neural Networks (CNNs), a powerful class of deep learning models that have revolutionized the field of computer vision. By the end of this week, you will have a solid understanding of CNNs and will be able to implement a basic CNN for image classification.

### Topics to be Covered:

- 1. Introduction to CNNs**
  - Understanding the basic structure and functionality of CNNs.
  - Differences between CNNs and traditional neural networks.
- 2. Convolutional Layers and Filters**
  - Learning about convolution operations and how filters (kernels) work.
  - Exploring how convolutional layers extract features from images.
- 3. Pooling Layers and Feature Extraction**
  - Understanding the purpose of pooling layers (e.g., max pooling, average pooling).
  - How pooling helps in downsampling and retaining essential features.
- 4. Implementing a Basic CNN for Image Classification**
  - A step-by-step guide to building a simple CNN using TensorFlow and Keras.
  - Training the model on a dataset and evaluating its performance.

### Resources

To help you understand and implement the concepts covered this week, here are some recommended resources:

-  Day 5-Understanding CNN & Implementation | Live Deep Learning Community S...
- [Understanding Convolutional Neural Networks for NLP · Denny's Blog](#)
-  Convolutional Neural Network Tutorial In TensorFlow / Keras

## Assignment-2: Basic Classification Task using CNN (Deadline: 10th June)

### Problem Statement:

Develop a CNN model for object detection using TensorFlow Keras. Utilize images from the COCO dataset and construct a CNN architecture suitable for object detection. You may choose to start with a pre-trained model such as ResNet or VGG. The objective is to build a model that can accurately identify objects within images and evaluate its performance using standard metrics.

## Tasks:

### 1. Data Preparation:

- Download a subset of the COCO dataset containing a variety of object classes.
- Preprocess the images and annotations to a suitable format for training a CNN.

### 2. Model Development:

- Construct a CNN architecture from scratch or fine-tune a pre-trained model (e.g., ResNet, VGG).
- Ensure your model includes essential components such as convolutional layers, pooling layers, and fully connected layers for classification.

### 3. Training the Model:

- Train your CNN on the prepared dataset.
- Use techniques like data augmentation and regularization to improve model performance.

### 4. Evaluation:

- Evaluate your model on a separate test set not used during training.
- Report metrics such as precision, recall, F1 score, and any other relevant performance metrics.

### 5. Reporting:

- Document your model architecture, training process, and results.
- Provide a detailed analysis of the model's performance and potential improvements.

By the end of this checkpoint, you should have a functional CNN model capable of performing basic object detection. Good luck!

## Submission Instructions

Submit the code and analysis on your respective GitHub and Google Drive link as stated in the [Submission Form](#).

## Extra for Interested Learners

For those of you who are particularly interested in the inner workings of CNNs, a fascinating exercise is to visualize how kernels evolve during training. By printing the images of kernels over time, you can see what features these kernels are learning to detect.